



Decision Tree and Random Forest



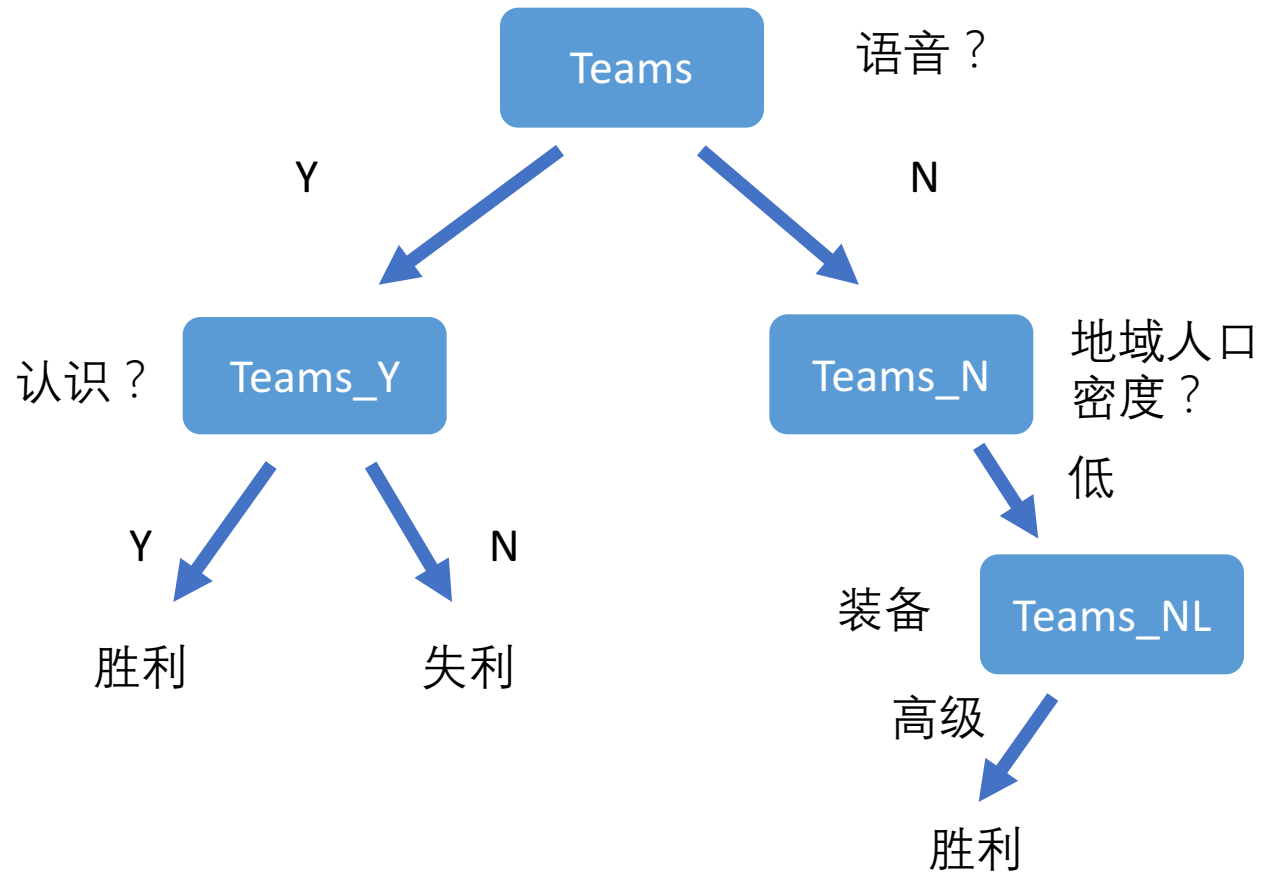
一个神奇的游戏：

我们小时候，可能都玩过这么一个游戏：

问一个问题，回答 是 或者 否 然后来猜你心中想的那个人



What is a Tree?



[Picture retrieved from www.dayinhu.com](http://www.dayinhu.com)





Take a look at trees

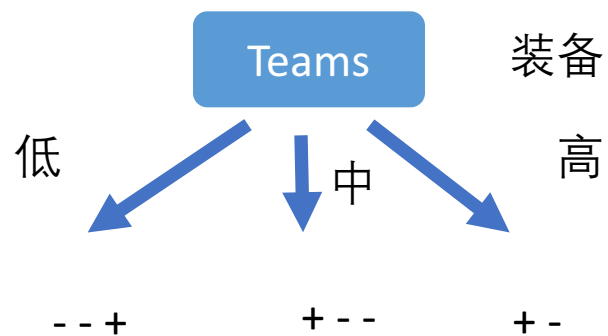
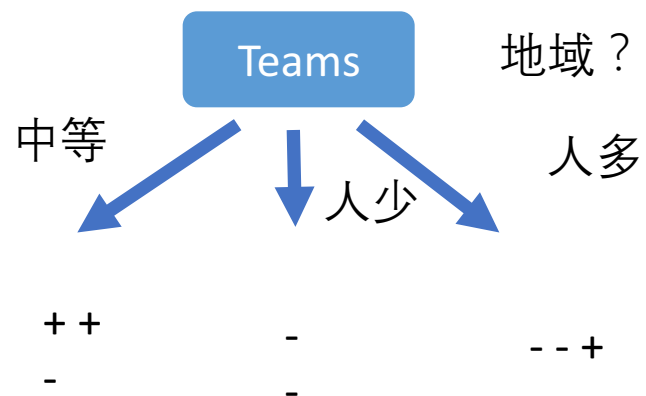
让我们开始问问题！

Win_the_game	know_each_other	Talk	Population	equipment
No	Not sure	Yes	few	low_level
No	Yes	Yes	many	low_level
Yes	Not sure	No	many	low_level
Yes	No	No	average	high_lavel
Yes	Not sure	No	average	medium_level
No	Yes	No	few	high_lavel
No	Yes	No	average	high_lavel
No	Not sure	Yes	many	medium_level

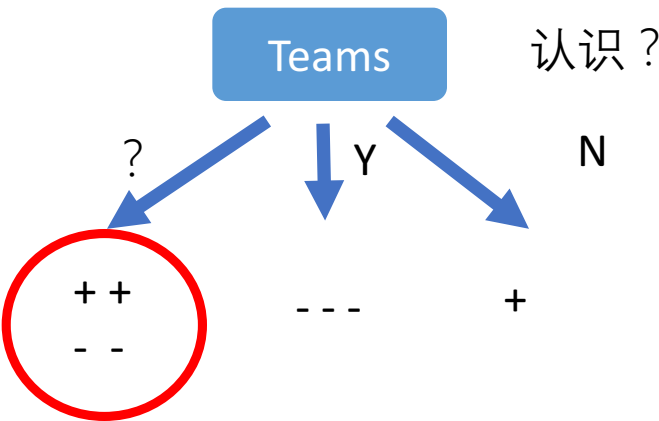
语音？






语音：3分
认识：4分
地域：2分
装备：0分



选了认识以后？



Win_the_game	know_each_other	Talk	Populati on	equipment
No	Not sure	Yes	few	low_level
No	Yes	Yes	many	low_level
Yes	Not sure	No	many	low_level
Yes	No	No	average	high_lavel
Yes	Not sure	No	average	medium_level
No	Yes	No	few	high_lavel
No	Yes	No	average	high_lavel
No	Not sure	Yes	many	medium_level





到底应该从哪个问题开始分？

到底哪个问题分的最好？

到底哪个问题分的结果最纯洁/纯粹？

什么叫纯粹？

$$Entropy = \sum_j -p_j * \log_2(p_j)$$

$$Gini = 1 - \sum_j p_j^2$$

From Information Theory

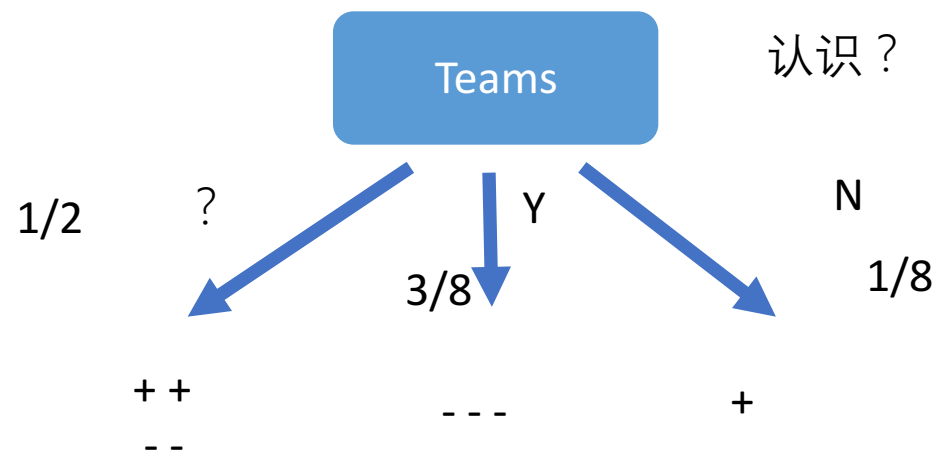
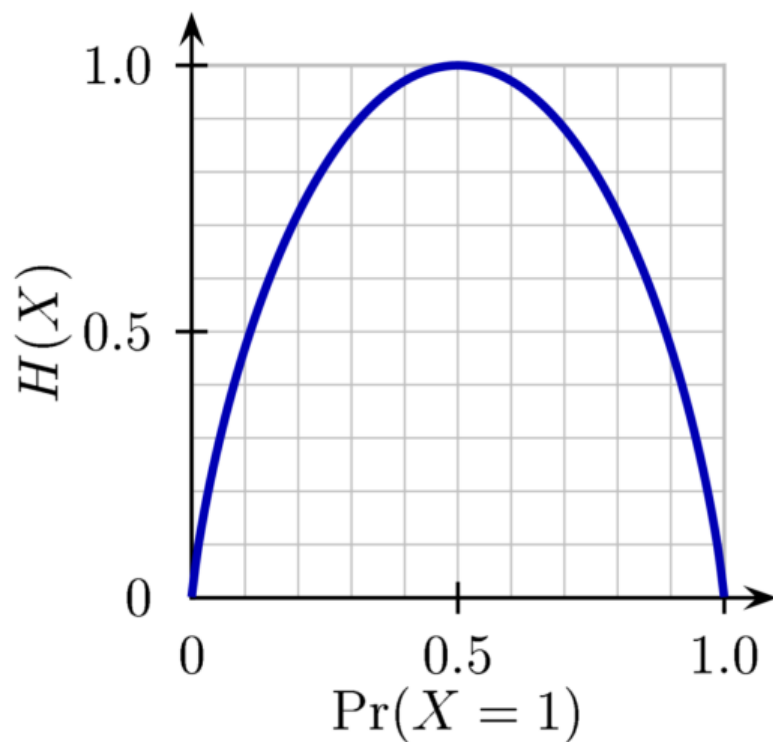
Gini Impurity

越低越好

$$Entropy = \sum_j -p_j * \log_2(p_j)$$



注意：
这里算的熵是对于**每一个叶子**，要想评价一个node，必须把**每一个叶子按权重比例加起来**



$$Entropy = 0.5 * 1 + 3/8 * 0 + 1/8 * 0 = 0.5$$

Picture retrieved from wiki : Binary entropy



How to deal with overfitting?

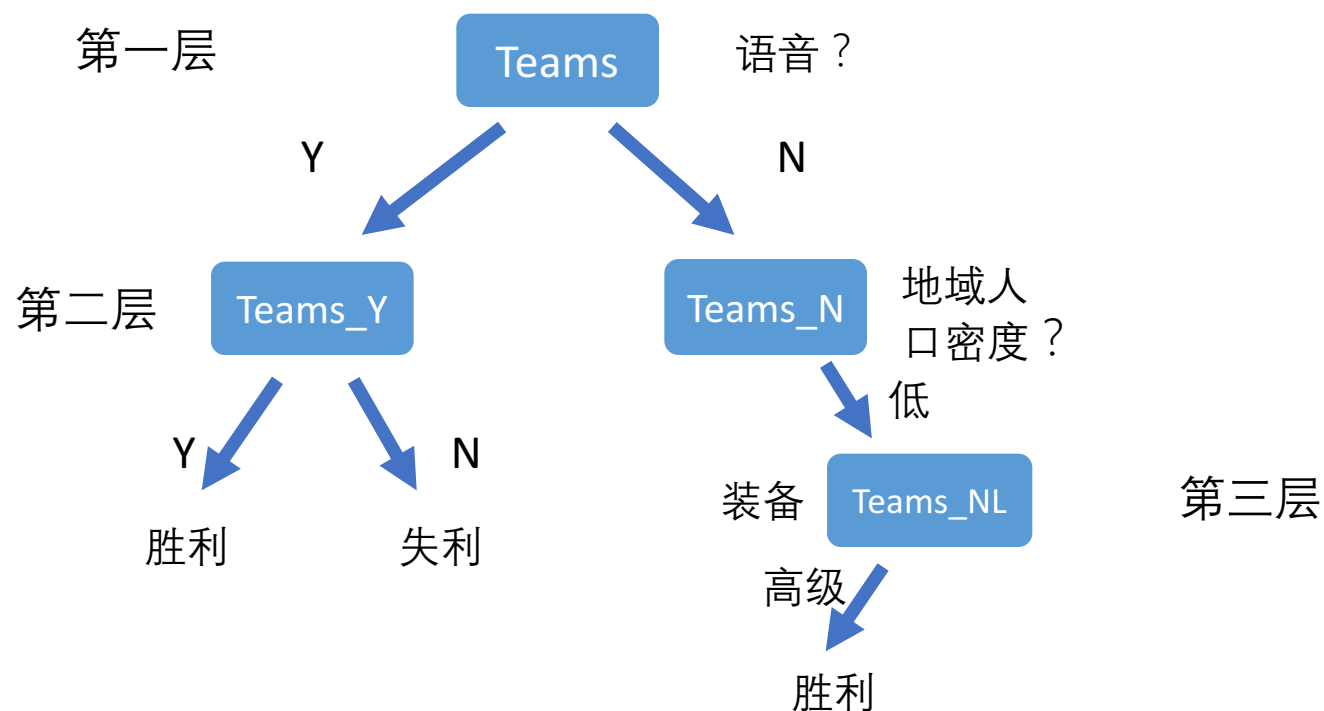
1. 限制能分几层 (number of split)

2. 最少可以分的数量

3. 每个叶子含有的样本数

4. 至少要增加多少information (information gain)

...



```
DTree=DecisionTreeClassifier(max_depth=5,criterion='gini',min_samples_split=2)
```



怎么把 Decision Tree 变的更强大 ？

- Bagging (Breiman, 1996) Bagging
- Boosting (Freund & Schapire, 1996) Boosting
- AdaBoost (Freund & Schapire, 1997) Boosting
- Random Forrest (Breiman, 1999) Bagging
- Gradient Boosting (Friedman et al, 2000) Boosting

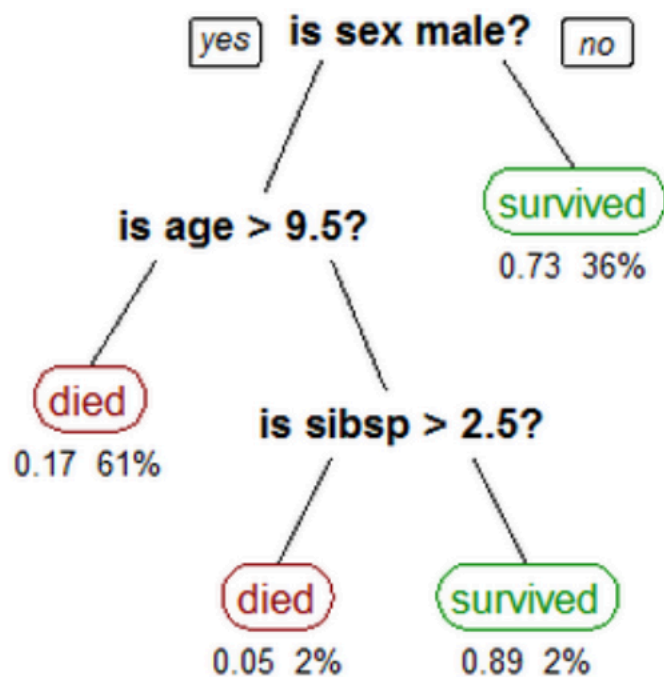
这些方法就叫做 集成算法 (Ensemble Learning)

Regression Tree 回归树



Outputs a number, 意思就是结果出来的是一个数

对比之前 Classification Tree 出来的是一个类别 (赢了 or 输了)



Titanic Survival Example

Bagging (Bootstrap aggregating)



假设有10个样本：

$[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]$

有放回的取10个样本：

$[2, 3, 1, 5, 2, 5, 6, 7, 7, 4]$

$[1, 1, 7, 0, 2, 4, 5, 5, 2, 9]$

...



Bagging (Bootstrap aggregating)

当样本量足够大时，取出的样本数量接近于63.2%

$$P = 1 - \left(1 - \frac{1}{N}\right)^N$$

If $N = 10$ $p = 65.1\%$

If $N = 100$ $p = 63.2\%$

好处：

- 模型更稳定
- 防止过拟合
- 提高准确率



Bagging (Bootstrap aggregating)

如何决定结果：少数服从多数





What about Random Forrest?

(随机 森林) = 随机 + 树 + 树 + 树...

- 用Bagging 来选取样本
- 用Bagging 来选取特征值



Boosting

本意：助推



“*... an efficient algorithm for converting **relatively poor** hypotheses into **very good hypotheses**...*”

能不能联合一帮弱的
变成一个强的？

— [Thoughts on Hypothesis Boosting](#) [PDF], 1988

Ada Boosting was invented in 1995

“*Boosting refers to this general problem of producing a very accurate prediction rule by combining rough and moderately inaccurate rules-of-thumb.*

— [A decision-theoretic generalization of on-line learning and an application to boosting](#) [PDF], 1995

Ada Boosting.M1



1. Initialize the observation weights $w_i = 1/N$, $i = 1, 2, \dots, N$.
2. For $m = 1$ to M :
 - (a) Fit a classifier $G_m(x)$ to the training data using weights w_i .
 - (b) Compute
$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$
 - (c) Compute $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$.
 - (d) Set $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$, $i = 1, 2, \dots, N$.

3. Output $G(x) = \text{sign} \left[\sum_{m=1}^M \alpha_m G_m(x) \right]$.

依次列出m 个分类器：

- 每次对于错误的样本，增加权重（起始的时候权重都是 $1/N$ ） w_i
- 对于预测的好的分类器（err 小）给予更多的权重 α_m
- 以少数服从多数的原则进行分类
- 最初是二分类 $\{-1, 1\}$

Ada Boosting 更深层次的理解



其实 Ada Boosting 所做的就是在做 Forward Stage Learning, 然后选取了 exponential 作为 Loss Function

$$Y_i = f_i(x) + \beta * b(xi; r)$$

每一步都找一个模型 $b(xi; r)$
去 fit 和真实值之间的差距

Exponential Loss Function :

$$L(y, f(x)) = \exp(-yf(x))$$

Ada Boosting 就是相当于用了
Exponential 作为 Loss Function

$$\alpha = 2\beta$$

Gradient Boosting

Ada Boosting 的一种通用形式



- Invent Adaboost, the first successful boosting algorithm
[Freund et al., 1996, Freund and Schapire, 1997]
- Formulate Adaboost as gradient descent with a special loss function
[Breiman et al., 1998, Breiman, 1999]
- Generalize Adaboost to Gradient Boosting in order to handle a variety of loss functions
[Friedman et al., 2000, Friedman, 2001]

Gradient Boosting



Consider you have series of predictions of x : $F(x_1)$, $F(x_2)$, $F(x_3)$, $F(x_4)$, ... $F(x_N)$,

But they are not quite accurate when comparing to y

For Example:

$$y_1=1.5 \quad \text{but } F(x_1) =1.34$$

$$y_2=0.8 \quad \text{but } F(x_2) =0.82$$

$$y_1=2.0 \quad \text{but } F(x_1) =1.91$$

**You may choose to
ignore it**

**Or fit the left off
with some model**

Gradient Boosting



Say if I use a model h to fit the difference between y and $F(x)$:

Use $h(x_1)$ to fit $y_1 - F(x_1)$
 $h(x_2)$ to fit $y_2 - F(x_2)$

这个叫残差

...

更新 $F'(x) = F(x) + h(x)$

但是如果 $F'(x)$ 还是不是完美的预测怎么办？

继续做一个模型 $b(x)$ 去 fit 新的残差

$y_1 - F(x_1) - h(x_1)$

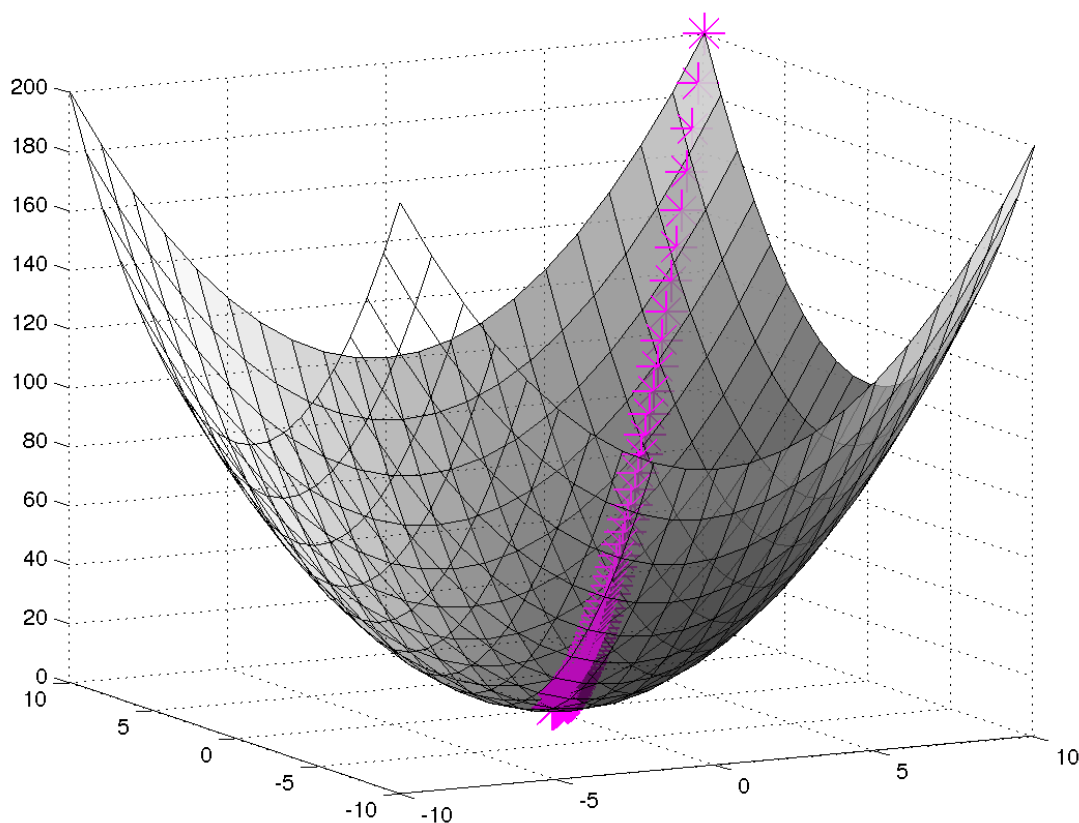
这就是Boosting 的哲学

为什么叫 Gradient Boosting ?



1. 大家都听说过 梯度下降 (gradient decent) 吧 :

怎么找到最小的位置 : 往 gradient 的反方向一小步一小步移动



$$\theta_i = \theta_i - \alpha \frac{\partial J}{\partial \theta_i}$$

J 在这里就是 Loss Function



如果我们用最小方差 作为 Loss Function, 或者说:

$$J = \frac{1}{2} \sum (y_i - F(x_i))^2$$

$$\frac{\partial J}{\partial F(x_i)} = F(x_i) - y_i$$

$$-\frac{\partial J}{\partial F(x_i)} = y_i - F(x_i)$$

Negative Gradient 相当于 残差

那么每次模型去 fit 残差 就是 模型去 fit Negative Gradient

每次根据残差去更新 $F'(x)$ 就是 根据 Negative Gradient 去更新 $F'(x)$

其实是用了 Gradient Descent 的思想 所以叫 Gradient Boosting (Gradient + Boosting)

当然 Gradient Boosting 所用的 Loss Function 不至于最小方差, 比如可以是 $\exp(-yf(x))$

A lush green forest scene with sunlight filtering through the trees. The sunbeams create a warm, golden glow in the center of the image. The forest floor is covered in green grass and fallen leaves. The trees are tall and slender, with dense foliage. The overall atmosphere is peaceful and serene.

谢谢观看！